

The Relational Model and SQL DDL

CISC637, Lecture #2


Ben Carterette

Copyright © Ben Carterette

1

Relational Model

- The relational model is the most widely used in modern DBMS
 - IBM's DB2, Informix, Oracle, Sybase, Microsoft Access and SQL Server, MySQL, PostgreSQL
- Introduced by Edgar Codd in 1970
- A set-based logical model for representing data



Copyright © Ben Carterette

Relations

- **Relations** are the main construct in the relational model
- Two components:
 - **Relation instance** is a table with rows and columns
 - A **relation schema** describes the columns of the table

Copyright © Ben Carterette

3

Relation Schema

- The schema defines a relation in terms of:
 - The relation name
 - The name of each column (or field)
 - The domain of each field
- Example:
 - `Student(ID: integer, name: string, dept_name: string, tot_cred: integer)`

Copyright © Ben Carterette

4

Relation Instance

- An **relation instance** is a set of records or tuples
 - Each record has the same fields as defined in the schema
 - The value in each field is in the field's domain as defined in the schema
 - Or a NULL value (allowed for some fields)
 - Every record is unique
- If any of these conditions are not true, then the relation instance is **invalid**

Copyright © Ben Carterette

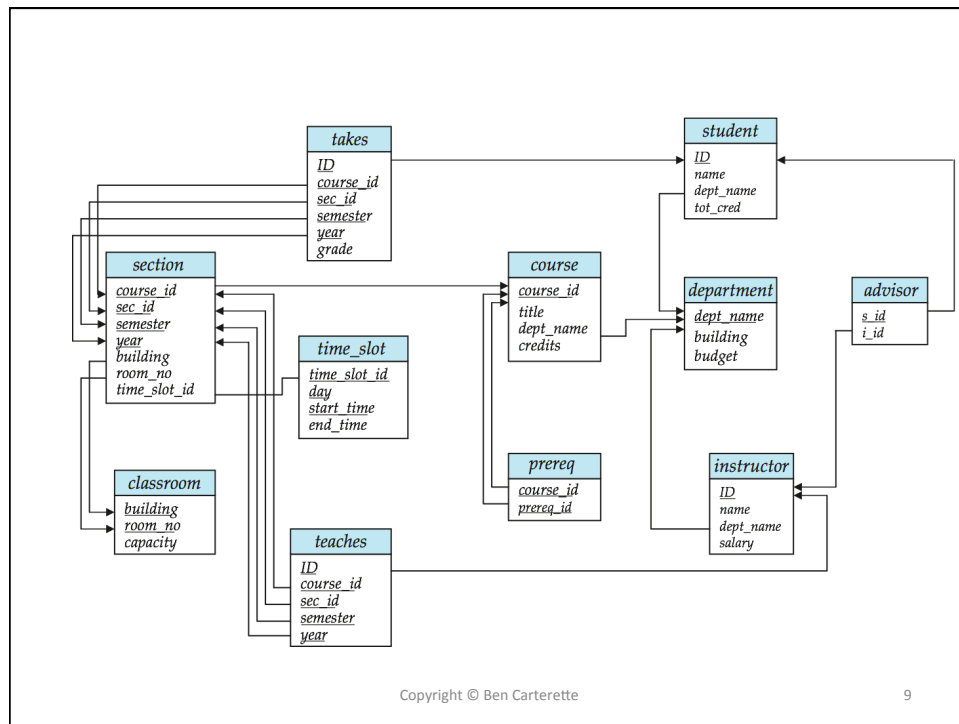
5

Relational Databases

- A **relational database** is a collection of relations with unique names
- The **relational database schema** is the collection of schema for its relations
- An **instance** of the RDB is a collection of relation instances

Copyright © Ben Carterette

8



Integrity Constraints

- **Integrity constraints** are used to ensure a database instance is valid
 - Defined on the database schema
 - Restrict the data that can be stored in the database instance
- A DBMS enforces the integrity constraints by only ever storing legal instances
- Example: *domain constraints*
 - Values in a record's field must correspond to domain as defined in the schema
 - No DBMS will store any data that violates a domain constraint (MySQL just typecasts everything)

Key Constraints

- A **key** is a field(s) that uniquely identifies a record
 - A **key constraint** says that every relation must have at least one key
 - It is an integrity constraint that ensures that all records are uniquely identifiable
- Three types of keys:
 - **Candidate key:**
 - Any set of fields that can uniquely identify records
 - No two distinct records can have the same values in all of a candidate key's fields,
 - but any strict subset of fields *can* have the same values in two different records.
 - **Superkey:**
 - Any set of fields that *contains* a candidate key
 - **Primary key:**
 - One of the candidate keys is chosen to be the primary key
 - Often the “smallest” candidate key, just one or a few fields

Copyright © Ben Carterette

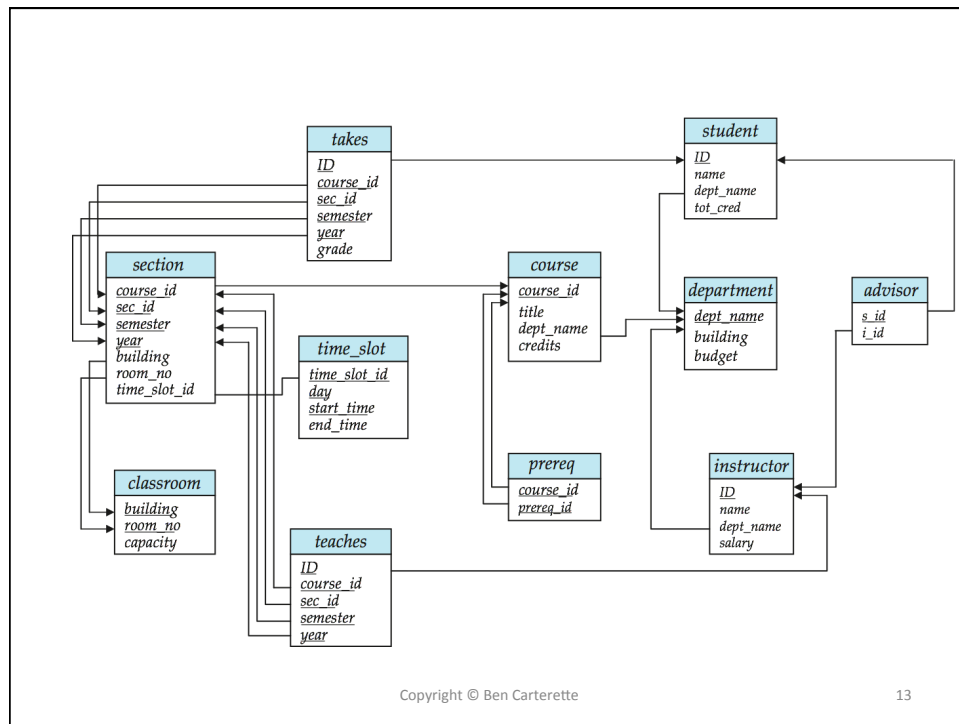
11

Foreign Key Constraints

- A **foreign key** is a set of fields in one relation that refers to a record in a *different* relation
 - In other words, a foreign key uniquely identifies a record in a different relation
- The purpose of foreign keys is to enforce **referential integrity**
 - Every foreign key should point to data that exists and is valid
 - No “dangling references”
- A DBMS that *enforces* foreign key constraints *guarantees* referential integrity

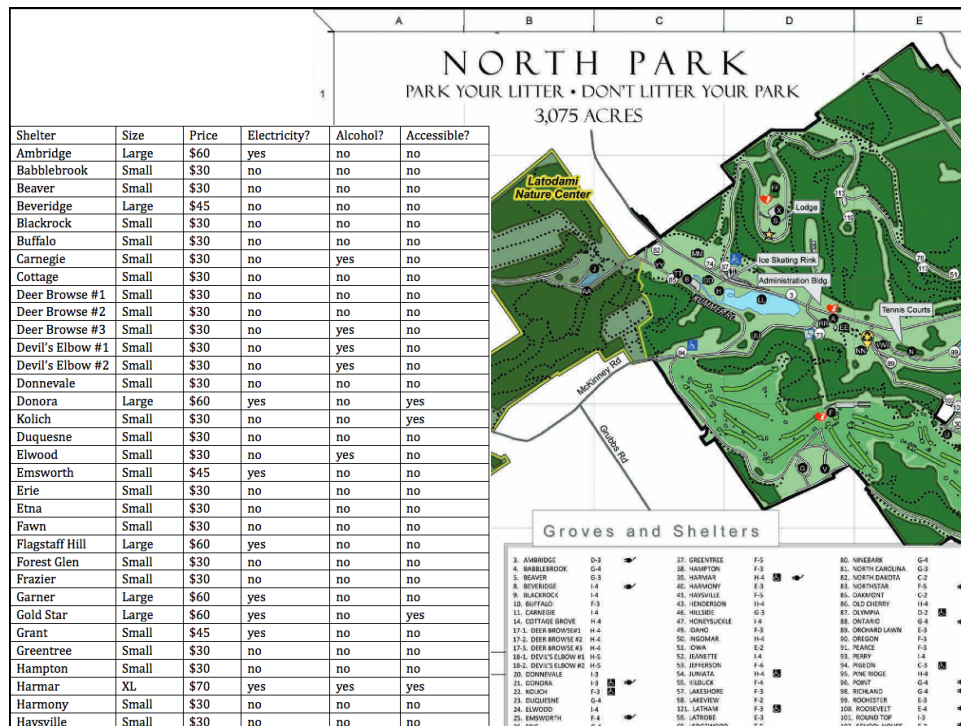
Copyright © Ben Carterette

12



Activity

- Same problem as Tuesday
 - Reserving a shelter at a public park
- This time, design the relational database
 - Write out tables with fields
 - Underline primary keys
 - Use arrows to indicate foreign keys
- Be sure to include *all* tables that would be necessary to use the app



SQL: A Brief Overview

- SQL: Structured Query Language
- Developed by IBM in the 1970s (separately from Codd)
- Standard language for DBMS
 - SQL:2011 is the current standard
 - Many DBMS include their own additions and modifications to the SQL language
- SQL comprises:
 - a Data Definition Language (DDL) for defining relations
 - a Data Manipulation Language (DML) for querying relations
- Today: a brief look at the DDL

Creating Relations in SQL

- Relations are always called “tables” in SQL

```
CREATE TABLE Student ( ID          INTEGER AUTO_INCREMENT,  
                        name        CHAR(30),  
                        dept_name   CHAR(40),  
                        tot_cred   INTEGER,  
                        PRIMARY KEY (ID) )
```

Copyright © Ben Carterette

17

Defining Other Candidate Keys

- Use the UNIQUE keyword to define other sets of fields that can uniquely identify records


```
CREATE TABLE Student ( ID          INTEGER AUTO_INCREMENT,  
                        ssn         INTEGER NOT NULL,  
                        name        CHAR(30),  
                        dept_name   CHAR(40),  
                        tot_cred   INTEGER,  
                        PRIMARY KEY (ID),  
                        UNIQUE (ssn) )
```

Copyright © Ben Carterette

18

Defining a Foreign Key

```
CREATE TABLE Student ( ID      INTEGER AUTO_INCREMENT,
                        name     CHAR(30),
                        dept_name CHAR(40),
                        tot_cred  INTEGER,
                        PRIMARY KEY (ID),
                        FOREIGN KEY (dept_name)
                        REFERENCES Department (dept_name) )
engine = InnoDB
```



This is required in MySQL to enforce foreign key constraints.
I am going to require that you always use it when defining tables in this class.

Copyright © Ben Carterette

19

Defining Foreign Keys

```
CREATE TABLE Takes (
  ID INTEGER,
  course_id VARCHAR(7),
  sec_id VARCHAR(3),
  semester VARCHAR(6),
  year INTEGER,
  grade VARCHAR(3),
  PRIMARY KEY (ID, course_id, sec_id, semester, year),
  FOREIGN KEY (ID) REFERENCES Student (ID),
  FOREIGN KEY (course_id, sec_id, semester_year)
  REFERENCES Section (course_id, sec_id, semester, year) )
engine = InnoDB
```

Copyright © Ben Carterette

20

Adding Records to a Table

- Insert a record into an existing table

```
INSERT  
INTO Student (ID, name, dept_name, tot_cred)  
VALUES (53688, 'Smith', 'CIS', 18)  
  
INSERT  
INTO Takes (ID, course_id, sec_id, semester, year, grade)  
VALUES (53688, 'CISC637', '010', 'Spring', '2015', 'A')
```

Copyright © Ben Carterette

21

Display All Records in a Table

- Very simple SQL query

```
SELECT *  
FROM Student
```

Copyright © Ben Carterette

22

Deleting Records in SQL

- Delete records that satisfy some logical condition

```
DELETE  
FROM Student  
WHERE name = 'Smith'
```

- But what should happen if there are records in Takes referring to students named Smith?

Copyright © Ben Carterette

23

Deleting Relations in SQL

- Deleting a table deletes schema as well as all records in that table

```
DROP TABLE Student
```

Copyright © Ben Carterette

24

Modifying Relations in SQL

- Rename tables; add/delete/rename columns; modify field domains

```
ALTER TABLE Student ADD birthdate CHAR(10)
ALTER TABLE Student DROP tot_cred
ALTER TABLE Student MODIFY birthdate DATE
ALTER TABLE Student ADD age INT
ALTER TABLE Student RENAME RegisteredStudent
```

Copyright © Ben Carterette

25

Modifying Records in SQL

- Change specified field values in all records that satisfy some logical condition

```
UPDATE Student
  SET tot_cred = tot_cred+3,
      birthdate = '1995-04-08'
WHERE ID = 53688
```

Copyright © Ben Carterette

26

Referential Integrity

- What should happen if:
 - ... we delete a record in Student that is referenced by a record in Takes?
 - ... we change the student ID of a Student that is referenced in Takes?
 - ... we delete a record in Course that is referenced in in a Section that is referenced in Takes?
 - ... we change the course ID of a Course that is referenced in a Section that is referenced in Takes?

Copyright © Ben Carterette

27

Enforcing Referential Integrity

- What to do when deleting or updating records that are referenced in other tables?
- Four options:
 - Delete/update all records that depend on the deleted record
 - Do not allow original record to be deleted/updated (throw error)
 - Set foreign key field values to refer to a default
 - Set foreign key field values to null

Copyright © Ben Carterette

28

Enforcing Ref. Integrity in SQL

```
CREATE TABLE Takes (
  ID INTEGER,
  course_id VARCHAR(7),
  sec_id VARCHAR(3),
  semester VARCHAR(6),
  year INTEGER,
  grade VARCHAR(3),
  PRIMARY KEY (ID, course_id, sec_id, semester, year),
  FOREIGN KEY (ID) REFERENCES Student (ID)
    ON DELETE CASCADE
    ON UPDATE CASCADE,
  FOREIGN KEY (course_id, sec_id, semester_year)
    REFERENCES Section (course_id, sec_id, semester, year)
    ON DELETE NO ACTION
    ON UPDATE NO ACTION )
engine = InnoDB
```

Copyright © Ben Carterette

29

Summary

- A *relational database* consists of *relations*
 - Relations are implemented as tables of records that have the same fields (but different values)
 - Each relation has a *relation instance* and a *relation schema*
- A valid relational database satisfies all integrity constraints defined on it
 - Domain constraints, key constraints, foreign key constraints
- The SQL Data Definition Language allows us to:
 - create relational databases
 - define integrity constraints
 - modify relation schema and data in relation instances

Copyright © Ben Carterette

30