

JOIN Syntax

- When you want to join tables in a chain, NATURAL JOIN can fail to understand how to complete joins

```
SELECT DISTINCT student.ID
FROM student S
      JOIN takes T ON S.ID=T.ID
      JOIN section C ON T.course_id=C.course_id
                        AND T.sec_id=C.sec_id
                        AND T.semester=C.semester
                        AND T.year=C.year
WHERE building='Smith';
```

Types of Joins

- INNER JOIN a.k.a. JOIN
 - Default join
 - Only returns records with values matching in both tables
 - NATURAL JOIN is a special type of INNER JOIN
- OUTER JOINS
 - LEFT OUTER JOIN
 - Result includes everything from INNER JOIN, plus records in left table that have no match in right table
 - RIGHT OUTER JOIN
 - Result includes everything from INNER JOIN, plus records in right table that have no match in left table
 - FULL OUTER JOIN
 - Result includes everything from INNER JOIN, plus records from both left and right tables that have no match in the other table
 - MySQL doesn't have this

Copyright © Ben Carterette

20

OUTER JOIN Example

- Find students who have not been assigned any advisor

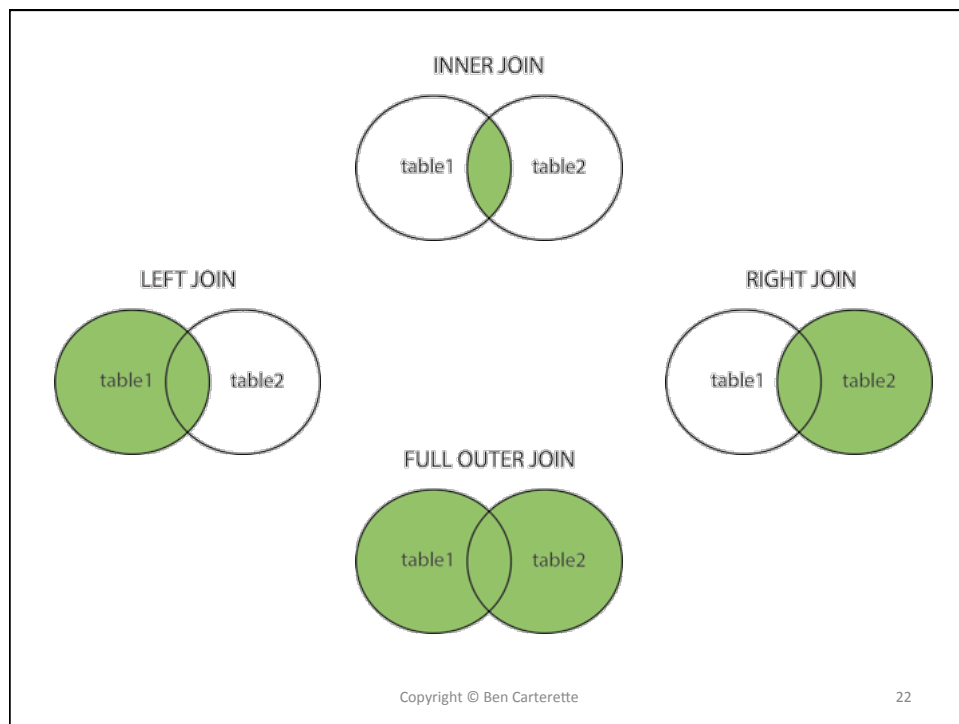
```
SELECT *
  FROM student S
      LEFT OUTER JOIN advisor A ON S.ID=A.s_id
 WHERE A.i_id IS NULL
```

- Compare to answer using nested subquery:

```
SELECT *
  FROM student S
 WHERE S.id NOT IN (SELECT s_id FROM advisor);
```

Copyright © Ben Carterette

21



Aggregation

```
SELECT field1, FUNCTION(field2), ...
FROM Table
WHERE condition
GROUP BY field1
```

- Calculate aggregate functions for values of field2 within *groups* of records by values of field1
 - AVG, MIN, MAX, SUM, COUNT are typical aggregation functions
- FUNCTION() computes something from values of field2
- GROUP BY restricts function to values of field2 in records with the same value of field1

Aggregation

- SELECT ... GROUP BY ...

<i>instructor</i>				<i>teaches</i>				
<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
22222	PHY-101			22222	PHY-101	1	Fall	2009

Copyright © Ben Carterette

24

More on Aggregation

- SELECT COUNT(DISTINCT field1) FROM Table
 - Return the number of unique values of field1
 - Different from SELECT DISTINCT COUNT(field1)!
- HAVING operator is like WHERE, but used to restrict by aggregated fields
 - SELECT field1, AVG(field2) AS mean FROM Table GROUP BY field1 HAVING mean > 3
 - Only return groups where the average value of field2 is more than three

Copyright © Ben Carterette

25

Aggregation

- SELECT .. GROUP BY .. HAVING ..

<i>instructor</i>				<i>teaches</i>				
<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>	<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>year</i>
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
22222				22222	PHY-101	1	Fall	2009

Departments with average monthly salary over \$10,000?

Semesters with more than 100 sections offered?

Copyright © Ben Carterette

26

Sorting

```
SELECT *
FROM Table1, Table2, ...
WHERE ...
ORDER BY Table1.field1
```

- Return results sorted in increasing order of Table1.field1
- Specify multiple fields to break ties
 - ORDER BY Table1.field1, Table2.field1, ...
- Use DESC or ASC to specify decreasing or increasing
 - ORDER BY Table1.field1 DESC, Table2.field1 ASC, ...
- Sorting can be expensive!

Copyright © Ben Carterette

27

Restricting Results by Row Number

```
SELECT *  
FROM Table1, Table2, ...  
WHERE Table1.fKey = Table2.pKey AND ...  
LIMIT n
```

- Returns only the first n matching records
- `LIMIT x, y` to select a range
 - `LIMIT 0, 10` equivalent to `LIMIT 10`
 - `LIMIT 5, 10` returns records 6, 7, 8, 9, and 10

Copyright © Ben Carterette

28

Conceptual Order of Operations

- First complete JOINS and WHEREs
 - If there is a chain of JOINS, MySQL decides the best way to complete them
 - Sometimes a WHERE is applied before a JOIN, other times the JOIN comes first—MySQL decides best order
- Then GROUP BY to form groups of those records with common field values
- Then HAVING to select a subset of groups
- Then ORDER BY to sort the result set
- Finally, if there is a DISTINCT or LIMIT, apply it last

Copyright © Ben Carterette

29

UNION

```
(SELECT field1, field2 FROM Table1 WHERE ...)  
UNION  
(SELECT field1, field2 FROM Table2 WHERE ...)
```

- The two tables must be **union-compatible**
 - Same number of fields, same domains
- Use UNION ALL to keep duplicates

Copyright © Ben Carterette

30

INTERSECT

```
(SELECT field1, field2 FROM Table1 WHERE ...)  
INTERSECT  
(SELECT field1, field2 FROM Table2 WHERE ...)
```

- The two tables must be union-compatible
 - Same number of fields, same domains
- Use INTERSECT ALL to keep duplicates
- MySQL does not support INTERSECT

Copyright © Ben Carterette

31

Set Intersection With Nested Subquery

- Implement INTERSECT using nested subqueries

```
SELECT course_id
FROM Teaches
WHERE semester='Spring' AND year=2015 AND
       course_id IN (SELECT course_id
                     FROM Teaches
                     WHERE semester='Fall' AND year=2014)
```

Copyright © Ben Carterette

32

EXCEPT (Set Difference)

```
(SELECT field1, field2 FROM Table1 WHERE ...)
EXCEPT
(SELECT field1, field2 FROM Table2 WHERE ...)
```

- The two tables must be union-compatible
 - Same number of fields, same domains
- Use EXCEPT ALL to keep duplicates
- MySQL does not support EXCEPT

Copyright © Ben Carterette

33

Set Difference With Nested Subquery

- Implement EXCEPT using nested subqueries

```
SELECT course_id
FROM Teaches
WHERE semester='Spring' AND year=2015 AND
      course_id NOT IN (SELECT course_id
                        FROM Teaches
                        WHERE semester='Fall' AND year=2014)
```

Copyright © Ben Carterette

34

Summary

- SQL is an extensive and deep query language for relational databases
 - Allows you to find records matching specific conditions across multiple tables
 - “Quickly find answers to arbitrary questions”
- Most important features:
 1. Joining tables
 2. Grouping rows and aggregating values
 3. Sorting
 4. Set operations

Copyright © Ben Carterette

35