

FDs and Normal Forms

CISC637, Lecture #10

Ben Carterette

Copyright © Ben Carterette

1

Database Design

- Database design is the process of going from *requirements* to *relational schema*
 - Schema should capture as many requirements as possible
 - Should also avoid redundancy and data inconsistency
 - Should also enable efficient storage and querying
 - These goals are not always mutually attainable
- There are many possible valid schema
 - One giant table that stores all data: valid design that can capture all requirements, but will have tons of redundancy
 - ...
 - Lots of small tables that store bits of data: valid design that has no redundancy but may not capture all requirements
- FDs and normal forms are tools for evaluating a possible design

Copyright © Ben Carterette

2

Design Example #1

- Change to university requirements:
 - Each instructor is an advisor for at most one department
 - Students can have majors in more than one department
 - Students can have multiple advisors, but at most one advisor per department
- FDs to capture these requirements:
 - $i_ID \rightarrow dept_name$
 - given instructor ID, there is at most one possible department they advise for
 - $s_ID, dept_name \rightarrow i_ID$
 - given student ID and department, there is at most one possible instructor advising

Copyright © Ben Carterette

3

Possible Relational Schema

- Possible relational schema #1:
 - Student(s_ID, name, tot_cred)
 - Instructor(i_ID, name, salary, dept_name)
 - DeptAdvisor(s_ID, dept_name, i_ID)
- Possible relational schema #2:
 - Student(s_ID, name, tot_cred)
 - StudentDept(s_ID, dept_name)
 - Instructor(i_ID, name, salary, dept_name)
 - Advisor(s_ID, i_ID)
- Both have pros and cons

Copyright © Ben Carterette

4

Design Example #2

- New university requirement
 - All students must have exactly one advisor
 - All instructors must advise exactly one student
 - Student and instructor must be in the same department
- FDs:
 - $s_ID \rightarrow name, tot_cred$
 - $i_ID \rightarrow name, salary$
 - $s_ID \rightarrow i_ID$
 - $i_ID \rightarrow s_ID$
 - $s_ID, i_ID \rightarrow dept$
- Two possible schema:
 - StudentInstructor(s_ID, s_name, tot_cred, i_ID, i_name, salary, dept)
 - StudentInstructor(i_ID, i_name, salary, s_ID, s_name, tot_cred, dept)
 - 3NF? Yes to both! BCNF? Yes to both! Good design? Probably not.

Copyright © Ben Carterette

5

FDs and Normal Forms

- FDs support a sort of logical system
 - Each one is a sentence that needs to be true on all possible database instances
 - (they can be checked for truth in the data)
 - There are *rules of inference* that can be applied to find FDs other than the ones specifically derived from requirements
- The system is **closed**
 - FDs inferred using rules are true, even if not specified
 - All FDs that are true can be derived using the rules
- The **closure** of a set of FDs is the set of all FDs that can be found by recursively applying inference rules
 - We will use F to refer to a given set of FDs
 - F^+ to refer to its closure

Copyright © Ben Carterette

6

Transitivity

- If $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow C$
- Example:
 - New university requirement
 - All students must have exactly one advisor
 - All instructors must advise exactly one student
 - Student and instructor in same department
 - FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$

 6. #3 and #2: $s_ID \rightarrow i_ID, \text{inst name, salary}$
 7. #4 and #1: $i_ID \rightarrow s_ID, \text{student name, tot_cred}$
 8. #6 and #7: $s_ID \rightarrow i_ID, \text{inst name, salary, student name, tot_cred}$

7

Augmentation

- If $A \rightarrow B$, then $\{A, C\} \rightarrow \{B, C\}$
- Example:
 - FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$

 6. augment #1: $s_ID, i_ID \rightarrow \text{student name, tot_cred, i_ID}$
 7. augment #2: $i_ID, s_ID \rightarrow \text{instructor name, salary s_ID}$

Copyright © Ben Carterette

8

Reflexivity & Armstrong's Axioms

- If A is a set of attributes and B is a subset of A, then $A \rightarrow B$
 - Trivial FDs
 - But a specific rule is necessary for the inference rules to be sound and complete
- Armstrong's axioms:
 1. Reflexivity
 2. Augmentation
 3. Transitivity
 - Applying these rules recursively to a set F will:
 - give every possible FD that is true (completeness)
 - will never give an FD that is not true (soundness)

Copyright © Ben Carterette

9

Armstrong's Axioms

- Armstrong's axioms:
 1. Reflexivity

If A is a set of fields and B is a subset of A, then $A \rightarrow B$
 2. Augmentation

If $A \rightarrow B$, then $\{A, C\} \rightarrow \{B, C\}$
 3. Transitivity

If $A \rightarrow B$ and $B \rightarrow C$, then $A \rightarrow C$
 - Applying these rules recursively to a set of FDs will:
 - give every possible FD that is true (completeness)
 - will never give an FD that is not true (soundness)

Copyright © Ben Carterette

10

Union

- If $A \rightarrow B$ and $A \rightarrow C$, then $A \rightarrow \{B, C\}$
 - Not one of Armstrong's rules, but follows from them
- Example:
 - FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$

 6. union of #1 and #3: $s_ID \rightarrow \text{student name, tot_cred, i_ID}$
 7. union of #2 and #4: $i_ID \rightarrow \text{inst name, salary, s_ID}$

Copyright © Ben Carterette

11

Decomposition

- If $A \rightarrow \{B, C\}$, then $A \rightarrow B$ and $A \rightarrow C$
 - Also not one of Armstrong's rules
- Example:
 - FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$

 6. decompose #1: $s_ID \rightarrow \text{student name, s_ID} \rightarrow \text{tot_cred}$

Copyright © Ben Carterette

12

FD Closure Example

- A *contract* is an agreement that a *supplier* will supply *qty* items of a certain *part* to a certain *project* associated with a certain *dept*
 - Contracts are identified by a unique ID
 - When a part is needed for a project, the entire quantity of that part is purchased with a single contract
 - A department purchases at most one part from a supplier
- F = the set of FDs from requirements:
 - contract \rightarrow supplier, project, dept, part, qty, value
 - project, part \rightarrow contract
 - supplier, dept \rightarrow part
- Schema (based on first candidate key FD):
 - Contracts(contract, supplier, project, dept, part, qty, value)
- Apply Armstrong's rules to get F^+

Copyright © Ben Carterette

13

BCNF/3NF Definitions Redux

- A relation R is in BCNF iff, for all $X \rightarrow A$ in F^+ ,
 - $A \subseteq X$ (a trivial FD), or
 - X is a superkey for R
- A relation R is in 3NF iff, for all $X \rightarrow A$ in F^+ ,
 - $A \subseteq X$ (a trivial FD), or
 - X is a superkey for R, or
 - A is part of some candidate key for R

Copyright © Ben Carterette

14

Attribute Closure

- Armstrong's inference rules are computationally inefficient
 - We could be applying them for a long time before we've completed F^+
- But we usually don't really care much about F^+ itself
 - What we care about is:
 - Is there some FD that is not a key? (violates BCNF)
 - Is there some FD that has fields on the right that are not part of any candidate key? (violates 3NF)
- The **attribute closure** of a field A is the set of all fields that are functionally determined by A after recursive inference
 - If $A \rightarrow$ all fields, then A is a superkey
 - If every field produces a superkey (or nothing), then BCNF

Copyright © Ben Carterette

15

Attribute Closure

- An efficient algorithm exists to compute attribute closure
- given a field (or set of fields) A and a set of functional dependencies F :

```

let  $A^+ = A$ 
do:
  for each FD  $f \in F$ :
    let  $B$  = left side of  $f$ 
    let  $C$  = right side of  $f$ 
    if  $B \subseteq A^+$ :
      let  $A^+ = A^+ \cup C$ 
until  $A^+$  is unchanged
  
```

Copyright © Ben Carterette

16

Student \leq Advisor \Rightarrow Instructor

- FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$
- Possible schema:
 - StudentInstructor(sid, sname, tot_cred, iid, iname, salary, dept)
 - Is this in BCNF?
- Is s_ID a superkey on StudentInstructor?
 - Start by putting s_ID in A^+
 - For each FD, if left side is in current version of A^+ , put right side in A^+
 - After first loop over FDs, $A^+ = \{s_ID, \text{student name, tot_cred, i_ID, dept}\}$
 - After second loop, $A^+ = \{s_ID, \text{student name, tot_cred, i_ID, dept, instructor name, salary}\}$
 - Since A^+ contains all fields, s_ID is a superkey
 - Therefore FDs 1 and 3 do not violate BCNF

Copyright © Ben Carterette

17

Contracts Example

- FDs:
 - $\text{contract} \rightarrow \text{supplier, project, dept, part, qty, value}$
 - $\text{project, part} \rightarrow \text{contract}$
 - $\text{supplier, dept} \rightarrow \text{part}$
- Is {project, part} a superkey?
 - Yes
- Is {supplier, dept} a superkey?
 - No—violates BCNF
 - Verify using attribute closure algorithm
- Candidate key {supplier, dept, project} cannot be found using attribute closure
 - But it can be verified to be a candidate key

Copyright © Ben Carterette

18

Student \leq Advisor \Rightarrow Instructor

- FDs:
 1. $s_ID \rightarrow \text{student name, tot_cred}$
 2. $i_ID \rightarrow \text{instructor name, salary}$
 3. $s_ID \rightarrow i_ID$
 4. $i_ID \rightarrow s_ID$
 5. $s_ID, i_ID \rightarrow \text{dept}$
- Alternate possible schema:
 - Student(s_ID, s_name, tot_cred)
 - Instructor(i_ID, i_name, salary)
 - Advisor(s_ID, i_ID, dept)
 - Are these in BCNF?
 - Yes: s_ID (from FD 1) is a superkey for Student
 - i_ID (from FD 2) is a superkey for Instructor
 - {s_ID, i_ID} (from FD 5) is a superkey for Advisor
 - Only problem is that we have to do extra work to make sure every instructor is an advisor and every student has an advisee
 - It happened automatically in the previous design

19

Formal Database Design Process

- Write out requirements as clearly as possible
- Draw an E-R diagram showing all important entities and relationships
- Use requirements and E-R diagram to list functional dependencies
- Convert E-R diagram to relational schema, preserving as many FDs as possible by keeping related attributes together
- Use attribute closure to test each relation for BCNF
 - If a relation is not in BCNF, either decompose it into BCNF relations ...
 - ... or test for 3NF