

## Analysis Table

	heap	sequential	seq + primary index	secondary index
scan	$O(B)$	$O(B)$	$O(B_i)$ index scan	$O(B_i)$ index scan
= search	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_2 B_i + m)$
<> search	$O(B)$	$O(\log_2 B + m)$	$O(\log_2 B_i + m)$	$O(\log_2 B_i + m)$
insert	$O(1)$	$O(B)$	$O(B_i)$	$O(B_i)$
delete	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_2 B_i + m)$

Note:

- This table reflects the time it takes to resort a file or index after insertion:
  - $O(B)$  for a sequential file,  $O(B_i)$  for an index
- Deletions do not necessarily require resorting


# Binary Search

Sequential index on name

[illegible]

## Binary Tree Search

- Binary search algorithm induces an ordered binary tree

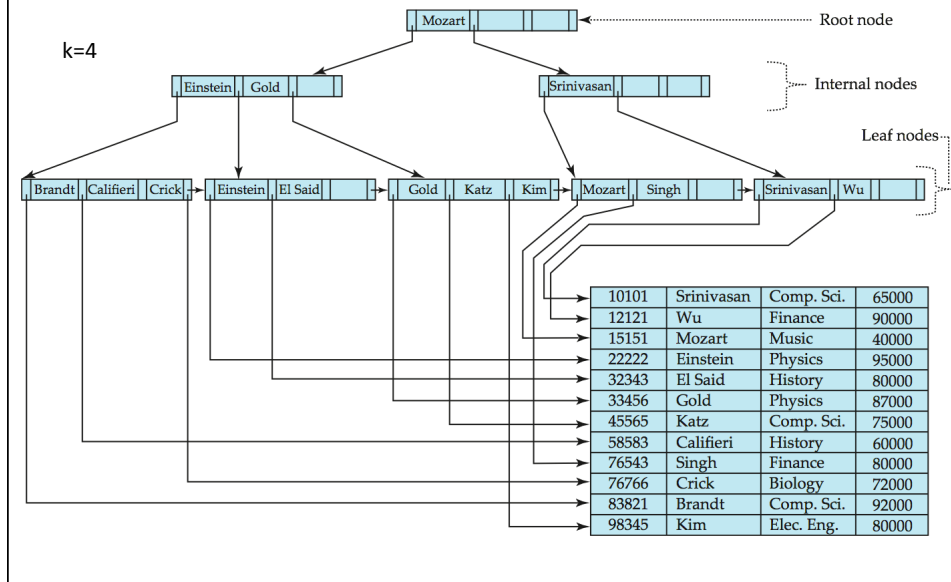


- What if we had a binary tree index instead of a sequential array index?


## B+ Tree Index (with a quick review of tree structures)

- An *ordered tree* is a tree in which children to the left of a parent have key values less than the parent; children to the right have key values greater
- A *balanced tree* is a tree in which, for any given parent node, the left subtree and right subtree are roughly the same height
  - Search, insert, delete in a balanced ordered tree is  $O(\log_2 n)$
- A *B-tree* is a balanced, ordered tree in which each node can have up to  $k$  children and  $k-1$  key values
  - A B-tree with  $k=2$  is a balanced binary tree
  - Search, insert, delete in a B-tree is  $O(\log_F n)$ , where  $F$  is the “fanout”
- A *B+ tree* is a B tree with key duplication
  - every key value in the index is stored in a leaf node
    - which means key values may be duplicated in internal nodes
- A **B+ tree index** is a multi-layer index
  - leaf nodes contain full index in sorted order
  - internal nodes are “indexes into indexes”
  - leaf nodes also store pointers to location of full records on disk

## Secondary B+ Tree Index on Name

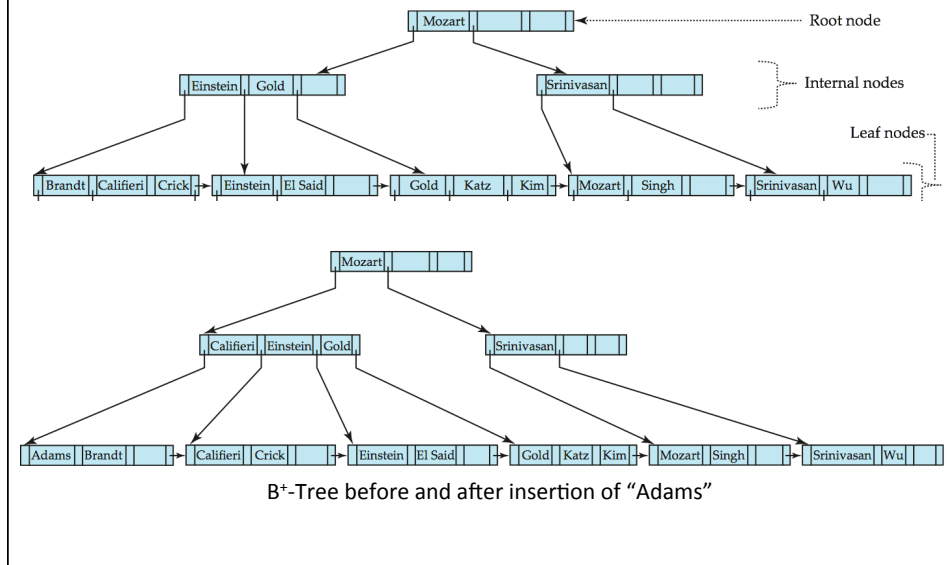


## Secondary B+ Tree Index on Name

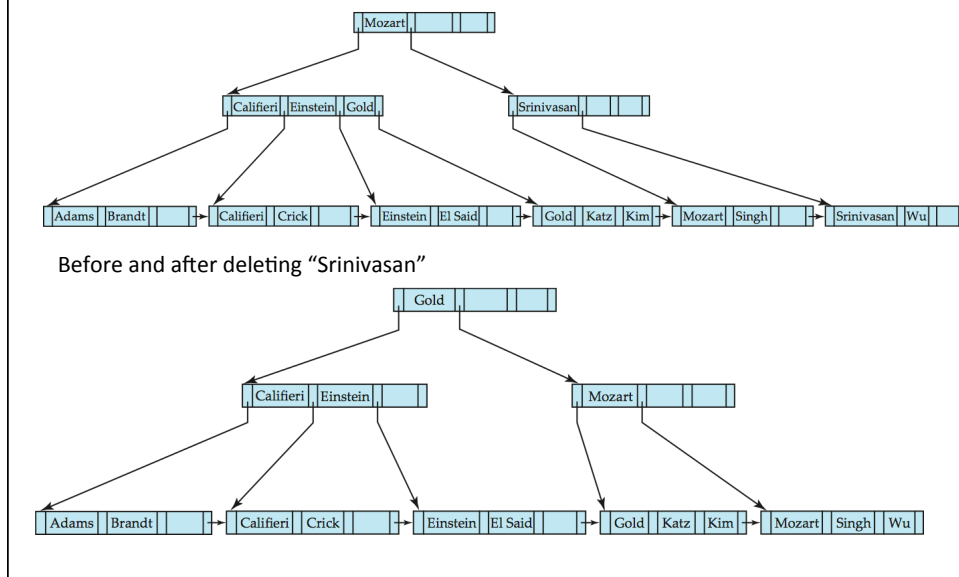


- Value of  $k$  influences height of tree
  - which in turn influences lookup/insert/delete time
  - ... and space overhead
- Index height =  $O(\log_{k/2} n)$ 
  - $n$  = total number of keys
  - Total blocks for storage in worst case = 1 (root node) +  $k/2$  (first layer) +  $k^2/4$  (second layer) + ... +  $B_l$  (leaf layer)
  - which works out to about  $O(B_l)$  extra blocks, or double the size of non-tree
- In real DBMS,  $k$  is very large (on the order of  $10^3$  or larger)
  - Results in a short, wide tree
  - Actual storage overhead much less than double size

## B+ Tree Insert




## B+ Tree Delete



## Sequential File With Primary B+ Tree Index

- Assume search key is primary key, so there is one index record per actual record
- Assume index records are 10% size of actual records
- Space costs:
  - Overhead?
- Time costs:
  - Scan?
  - Search with equality?
  - Search for range?
  - Insert?
  - Delete?




Sequential file with primary B+ tree index on ID (k=6)

## Sequential File With Primary B+ Tree Index

- Assume search key is primary key, so there is one index record per actual record
- Assume index records are 10% size of actual records
- Space costs:
  - Overhead?  $B_i + B_i/(k/2)$ 
    - $B_i$  typically  $0.15B$ , as blocks in B+ tree index are rarely 100% full
- Time costs:
  - Scan?  $B(D+RC)$  for full table scan,  $B_i(D+R_iC)$  for index scan
  - Search with equality?  $O(D \log_{k/2} B_i) + D$
  - Search for range?  $O(D \log_{k/2} B_i) + mD$
  - Insert?  $O(D \log_{k/2} B_i) + 4D$
  - Delete?  $O(D \log_{k/2} B_i) + 4D$

## Secondary B+ Tree Index

- File organization is a heap relative to the field being indexed
- Assume index records are 10% size of actual records
- Space costs:
  - Overhead?
- Time costs:
  - Scan?
  - Search with equality?
  - Search for range?
  - Insert?
  - Delete?



Imagine there is a B+ tree index on salary

## Analysis Table

	heap	sequential	seq + primary index	secondary index	primary B+-tree index	secondary B+-tree index
scan	$O(B)$	$O(B)$	$O(B_i)$ index scan	$O(B_i)$ index scan	$O(B_i)$ index scan	$O(B_i)$ index scan
= search	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_2 B_i + m)$	$O(\log_F B_i)$	$O(\log_F B_i + m)$
<> search	$O(B)$	$O(\log_2 B + m)$	$O(\log_2 B_i + m)$	$O(\log_2 B_i + m)$	$O(\log_F B_i + m)$	$O(\log_F B_i + m)$
insert	$O(1)$	$O(B)$	$O(B_i)$	$O(B_i)$	$O(\log_F B_i)$	$O(\log_F B_i)$
delete	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_2 B_i + m)$	$O(\log_F B_i)$	$O(\log_F B_i + m)$

Notes:

- Insert time for sequential files and non-tree indexes reflects time to re-sort:  $O(B)$
- For B+-trees, F is the fanout of the tree
  - $F = \text{ceiling}((k+1)/2)$

## Hash File Organization and Hash Indexes

- A **hash table** is a data structure mapping keys to values
  - Consists of a **hash function** and one or more **buckets**
  - Each bucket stores values associated with one or more keys
  - The hash function  $h(k)$  takes a key as input and outputs a bucket address
- All values associated with a given key will be put in the same bucket
- A single bucket may contain values for more than one key

## Hash File Organization

- Hash files organized so that all records with a given search key value are stored in the same block on disk
  - disk blocks are the buckets
  - $h(\text{search key}) = \text{block address}$
- Example: instructor relation organized as a hash file with dept\_name as search key
  - 10 buckets
  - $h(k)$  could be sum of integer values of characters mod 10
    - $h(\text{music}) = 1$ ,  $h(\text{history}) = 2$ ,  $h(\text{physics}) = 3$ ,  $h(\text{elec. eng}) = 3$ , etc

bucket 0


bucket 1

15151	Mozart	Music	40000

bucket 2

32343	El Said	History	80000
58583	Califieri	History	60000

bucket 3

22222	Einstein	Physics	95000
33456	Gold	Physics	87000
98345	Kim	Elec. Eng.	80000

bucket 4

12121	Wu	Finance	90000
76543	Singh	Finance	80000

bucket 5

76766	Crick	Biology	72000

bucket 6

10101	Srinivasan	Comp. Sci.	65000
45565	Katz	Comp. Sci.	75000
83821	Brandt	Comp. Sci.	92000

bucket 7

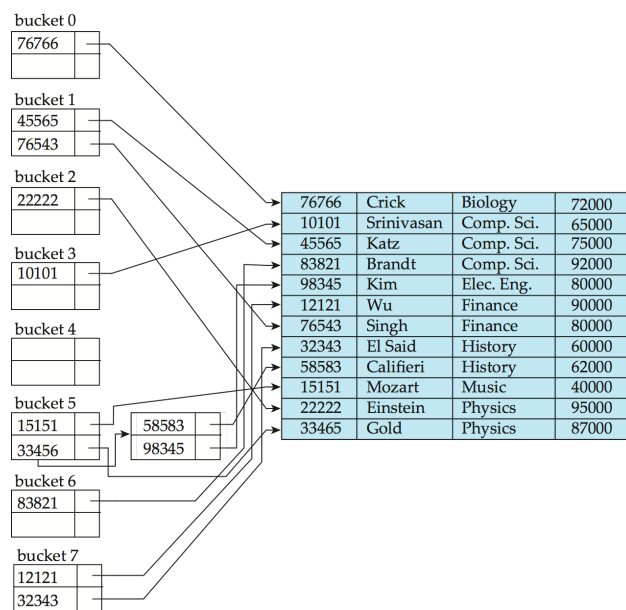

## Overflow

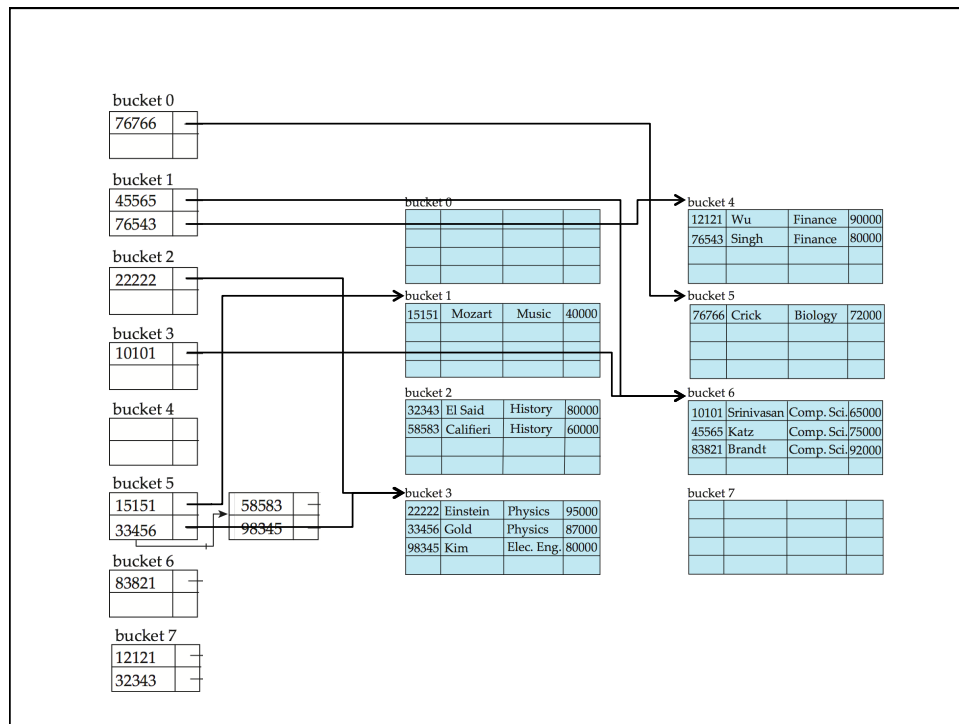
- Impossible to guarantee that each block/bucket will have enough space for all records that hash there
- **Overflow** buckets store records that don't fit in the original bucket
  - **Closed hashing:** overflow buckets chained together in linked list
  - Each block contains pointer to its first overflow block
    - Each overflow block contains pointer to the next one



## Hash Indexes

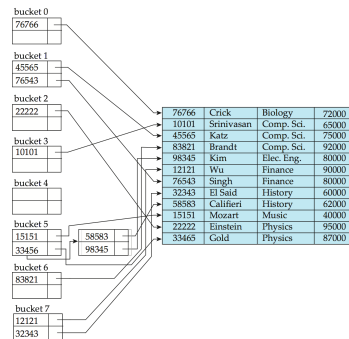
- A **hash index** maps search keys to blocks where index records are stored
  - Difference between hash index and hash file is that the index contains **only** search key values
- Note: a file organized using a hash function is already effectively indexed by a search key
  - Thus doesn't really make sense to talk about "primary indexes" on hash files
  - **All** hash indexes are secondary
  - But we will use "hash index" to refer to both a hash file org as well as a hash index





## Heap File With Hash Index

- Assume search key is primary key, so there is one index record per actual record
- Assume index records are 10% size of actual records
- Assume index blocks are at 80% capacity on average
- Space costs:
  - Overhead?
- Time costs:
  - Scan?
  - Search with equality?
  - Search for range?
  - Insert?
  - Delete?



Pretend this is a heap file

## Hash Indexes in Practice

- MySQL does not support hash indexes for tables stored on disk
  - It will let you execute `CREATE INDEX .. HASH`, but it will create a B+-tree index (without telling you)
  - Many DBMSs do not support them
- Why?
  - Difficult to predict how much space to reserve for index without knowing how many keys there will be
    - Which means it's difficult to decide on a good hash function
  - Efficiency of equality search with hash index rarely better than B+-tree with high  $k$
  - Completely useless for range search
  - Lose the nice side-effect of a B+-tree that key values are already sorted
    - Which often means more post-processing time to sort
  - All in all: time spent coding hash index generally not worth the small gains that can be achieved given the losses that might be incurred

## Analysis Table

	heap	sequential	seq + primary index	primary B+-tree index	secondary B+-tree index	hash
scan	$O(B)$	$O(B)$	$O(B_i)$ index scan	$O(B_i)$ index scan	$O(B_i)$ index scan	$O(B)$ or $O(B_i)$
= search	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_F B_i)$	$O(\log_F B_i + m)$	$O(B_m)$
<> search	$O(B)$	$O(\log_2 B + m)$	$O(\log_2 B_i + m)$	$O(\log_F B_i + m)$	$O(\log_F B_i + m)$	$O(B)$
insert	$O(1)$	$O(B)$	$O(B_i)$	$O(\log_F B_i)$	$O(\log_F B_i)$	$O(B_m)$
delete	$O(B)$	$O(\log_2 B)$	$O(\log_2 B_i)$	$O(\log_F B_i)$	$O(\log_F B_i + m)$	$O(B_m)$

Notes:

- Insert time for sequential files and non-tree indexes reflects time to re-sort:  $O(B)$
- For B+-trees,  $F$  is the fanout of the tree
  - $F = \text{ceiling}((k+1)/2)$
- For hash storage,  $B_m$  is the number of blocks in the relevant bucket
  - $B_m < B_i$
  - But difficult to quantify relationship between  $B_m$  and  $\log_F B_i$

## Indexes and Performance Tuning

- A **workload** is a mix of queries and update operations
- We'd like to create indexes that will support the expected workload efficiently
- For each query in the workload:
  - What relations does it access?
  - What attributes are retrieved?
  - Which attributes are involved in select/join clauses?
  - How selective are those conditions?
- For each update in the workload:
  - What type of update (INSERT/UPDATE/DELETE)?
  - What attributes are affected?

Copyright © Ben Carterette

50

## Choosing Indexes

- What indexes should we create?
  - Which relations need indexes?
  - What fields should be used as search keys?
  - Do we need more than one index for a relation?
- What type of index?
  - Clustered? Hash? B+-tree?
  - No choice in MySQL: all files stored as sequential files with B+-tree indexes on primary key, and all secondary indexes are B+-tree indexes

Copyright © Ben Carterette

51